

Predicting Diabetes with Machine Learning (Pima Indians Diabetes Dataset)

Name: Anurag Majumdar

Regd. No.: 2341013076

Sec: 2N3

Branch: B. Tech - CSE

Objective:

The goal of this project is to predict whether a patient has diabetes based on their medical details using Machine Learning.

This enables early detection and can assist healthcare professionals in decision-making.

Tools & Technologies Used

Tool / Library	Purpose
Python 3.8+	Programming language
Jupyter Notebook	Interactive environment
pandas, numpy	Data loading, cleaning, and processing
matplotlib, seaborn	Data visualization
scikit-learn	Machine learning model & evaluation

Installation:

```
pip install pandas numpy matplotlib seaborn scikit-learn
```

Dataset

- **Dataset:** [Pima Indians Diabetes Database](#)
- **Size:** 768 rows \times 9 columns
- **Target Column:** Outcome (1 = Diabetes, 0 = No Diabetes)

Feature	Description
Pregnancies	Number of times pregnant
Glucose	Plasma glucose concentration
BloodPressure	Diastolic blood pressure (mm Hg)
SkinThickness	Triceps skinfold thickness (mm)
Insulin	Serum insulin (mu U/ml)
BMI	Body Mass Index
DiabetesPedigreeFunction	Diabetes heredity score
Age	Age in years
Outcome	0 = No diabetes, 1 = Diabetes

Data Preprocessing

- Loaded data with pandas
- Checked for invalid/missing values
- Replaced 0 values in BMI and BloodPressure with median values
- Split features (X) and target (y)
- Standardized features using StandardScaler

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

Exploratory Data Analysis (EDA)

- **Histogram of Age & BMI** → Most patients were 20–40 years old, BMI skewed toward overweight range.
- **Outcome Distribution:**
 - **0 (No Diabetes):** 500+ patients
 - **1 (Diabetes):** ~260 patients

Visuals made patterns clearer — more non-diabetic patients, which could affect model balance.

Model Training

- **Split Data:** 70% training / 30% testing
- **Algorithm Used:** Logistic Regression
- **Code:**

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
y_pred = model.predict(X_test_scaled)
```

Confusion Matrix

Generated confusion matrix to evaluate predictions:

	Predicted No	Predicted Yes
Actual No	TN	FP
Actual Yes	FN	TP

- **TP (True Positive):** Diabetic patient correctly classified
- **TN (True Negative):** Healthy patient correctly classified
- **FP (False Positive):** Healthy patient predicted as diabetic (false alarm)
- **FN (False Negative):** Diabetic patient predicted as healthy (dangerous miss)

Key Insights

- Logistic Regression is a good **baseline model** for medical classification tasks.
- False negatives are more dangerous than false positives in healthcare — we must minimize FN even if it slightly increases FP.
- Scaling features improved accuracy.

Reflection

This project taught me how to take a dataset from raw form to a working prediction model step by step. I found data visualization and model training easy, but handling missing values and understanding scaling were a bit tricky at first. The biggest learning was interpreting the confusion matrix and realizing that in healthcare, minimizing false negatives is more important than just high accuracy. Overall, this project gave me confidence in working with real data and understanding the practical side of machine learning.